

MASARYK UNIVERSITY
FACULTY OF INFORMATICS



Use of Transactions within a Reactive Microservices Environment

MASTER'S THESIS

Martin Štefanko

Brno, Fall 2017

MASARYK UNIVERSITY
FACULTY OF INFORMATICS



Use of Transactions within a Reactive Microservices Environment

MASTER'S THESIS

Martin Štefanko

Brno, Fall 2017

Replace this page with a copy of the official signed thesis assignment and the copy of the Statement of an Author.

Declaration

Hereby I declare that this paper is my original authorial work, which I have worked out on my own. All sources, references, and literature used or excerpted during elaboration of this work are properly cited and listed in complete reference to the due source.

Martin Štefanko

Advisor: Bruno Rossi, PhD

Acknowledgement

thanks

Abstract

abstract

Keywords

transactions, Narayana, JTA, reactive, microservices, asynchronous, saga, compensating transactions

Contents

1	Introduction	1
2	Transaction concepts	2
2.1	<i>Transaction</i>	2
2.2	<i>2PC protocol</i>	2
2.3	<i>ACID properties</i>	2
2.3.1	Atomicity	3
2.3.2	Consistency	3
2.3.3	Isolation	3
2.3.3.1	Isolation levels	4
2.3.4	Durability	4
2.4	<i>Transaction models</i>	4
2.5	<i>Distributed transactions</i>	4
2.6	<i>Transaction manager</i>	4
2.6.1	Local transaction manager	5
2.7	<i>Failure handling</i>	5
3	Microservices architecture pattern	6
4	Communication patterns	7
4.1	<i>Consensus protocols</i>	7
4.1.1	2PC	7
4.1.2	3PC raft paxos	7
4.2	<i>Event based protocols</i>	7
4.2.1	CQRS	7
5	Transactions in distributed environment	8
5.1	<i>Saga pattern</i>	8
5.1.1	Participants	8
6	Conclusion	9
	Bibliography	10

1 Introduction

2 Transaction concepts

This chapter introduces the basic notions of transactions, their properties and common problems of the management of transactions across multiple nodes in the distributed systems.

2.1 Transaction

A transaction is an unit of processing that provides all-or-nothing property to the work that is conducted within its scope, also ensuring that shared resources are protected from multiple users [1]. It represents an unified and inseparable sequence of operations that are either all provided or none of them take effect.

From the application point of view there exist several transaction models in which the transactions can be executed. The applicable models in the transaction management are local, programmatic and declarative transaction models. All three models will be described in detail in the following section.

The transaction can end in two forms: it can be either *committed* or *aborted*. The commit determines the successful outcome - all operations within the transaction have been performed and their results are permanently stored in a durable storage. The abort means that all performed operations have been undone and the system is in the same state as if the transaction have not been started.

Generally the achieving of above features may differ. The most common pattern for the transaction processing is a two phase commit protocol with the ACID transactions. Other approaches are based on the relaxation of the one or more of ACID properties to adjust to the real world environments.

2.2 2PC protocol

2.3 ACID properties

A transaction can be viewed as a group of business logic statements with certain shared properties [2]. Generally considered properties are

one or more of atomicity, consistency, isolation and durability. These four properties are often referenced as ACID properties [3] and they describe the major points important for the transaction concepts.

2.3.1 Atomicity

The transaction consists of a sequence of operations performed on different resources or by different participants. Atomicity means that all operation in the transaction are performed as if they were a single operation. When the transaction commits successfully all of its participants are also required to perform a valid commit. Conversely, if the transaction fails and is aborted all performed operations and effects are forced to be undone. This defines a possibility to abort at any point so that all changes done by the transaction will be reverted to the state before the transaction start.

Atomicity is generally achieved by the usage of the consensus multi-phase protocols. Standardized protocol is the two phase commit protocol which is used by the majority of modern transaction systems.

2.3.2 Consistency

Consistency describes that the transaction maintains the consistency of the system and resources that it is performed on. When the transaction is started on the consistent system this system must remain consistent when the transaction ends - it moves from one consistent state to another.

Unlike other transactional properties (A, I, D), consistency cannot be realized by the transaction system as it does not hold any semantic knowledge about the resources it manipulates [1]. Therefore achieving this property is the responsibility of the application code.

2.3.3 Isolation

An isolation property takes effect when multiple transactions can be executed concurrently on the same resources. That means concurrent transactions can not interfere one with another. Therefore each concurrent execution on the shared resource must be equivalent to some

serial ordering of contained transactions which is why the isolation is often also referenced as serializability.

From the perspective of an external user the isolation property means that the transaction appears as it was executed entirely by itself. This means that even if there are multiple transactions in the system executed concurrently, this fact is hidden from the external views when the system is serializable.

As an instinctive extension of the consistency property, the serial execution of the transaction keeps the consistent state. The execution of the transactions in parallel therefore cannot result into the inconsistent system.

2.3.3.1 Isolation levels

2.3.4 Durability

2.4 Transaction models

<https://ress.infoq.com/minibooks/JTDS/en/pdf/javatransactionsbook.pdf?Expires=1512505329&Signature=NBftdwZSbBAKvAxQ1LfPMzv0~9vusCCd1WxWIqcSU7.&Key-Pair-Id=APKAIMZVI7QH4C5YKH6Q>

2.5 Distributed transactions

A distributed transaction is the transaction performed in a distributed system. The distributed system consists of a number of independent devices connected through a communication network. Such systems are liable to the frequent failures of individual participants or communication channels between them.

The transaction manager can be implemented as a separate service or being placed with some participant or the client. **TODO**

2.6 Transaction manager

Every transaction is associated with a transaction coordinator or transaction manager which is responsible for the control and supervision of the participants performing individual operations. It is a component liable for coordinating transactions in the sequential or parallel

execution across one or more resources. It provides proper and complete execution and it administers the comprehensive result of the transaction. Applications are commonly required only to contact the transaction manager about the start of the transaction.

The main responsibilities of transaction manager are starting and ending (commit or abort) of the transaction, management of the transaction context, supervision of transactions scoped across multiple resources and the recovery from failure.

2.6.1 Local transaction manager

A local transaction manager or a resource manager is responsible for the coordination of transactions concerning only a single resource. Because of its scope it is often build in directly to the resource. The span of the resource is defined by its managing platform.

The resource manager is required to provide a support for the participation in the global transactions that span over several resources. This means that it is effectively capable to handle complete transaction processing to the different transaction manager.

2.7 Failure handling

CA CP theorem - CAP

3 Microservices architecture pattern

4 Communication patterns

4.1 Consensus protocols

4.1.1 2PC

4.1.2 3PC raft paxos

4.2 Event based protocols

Eventual consistency

4.2.1 CQRS

5 Transactions in distributed environment

5.1 Saga pattern

A saga, as described in the original publication [4], is a long lived transaction that can be written as a sequence of transactions that can be interleaved with other transactions. Each operation that is a part of the saga represents an unit of work that can be undone by the compensation action. The saga guarantees that either all operations complete successfully, or the corresponding compensation actions are run for all executed operations to cancel the partial processing.

5.1.1 Participants

6 Conclusion

Bibliography

- [1] M. Little, J. Maron, and G. Pavlik, *Java transaction processing*. Prentice Hall, 2004.
- [2] M. Musgrove, “Narayana + wildfly,” 2015. [Online]. Available: <https://developer.jboss.org/servlet/JiveServlet/download/53044-3-129391/jbug-brno-transactions.pdf>
- [3] T. Haerder and A. Reuter, “Principles of transaction-oriented database recovery,” *ACM Computing Surveys*, vol. 15, no. 4, pp. 287–317, 1983.
- [4] H. Garcia-Molina and K. Salem, “Sagas,” *ACM SIGMOD Record*, vol. 16, no. 3, pp. 249–259, 1987.